

Data Structures & Algorithms

Lecture 5: Hash Tables

Chapter 11

Abstract Data Types

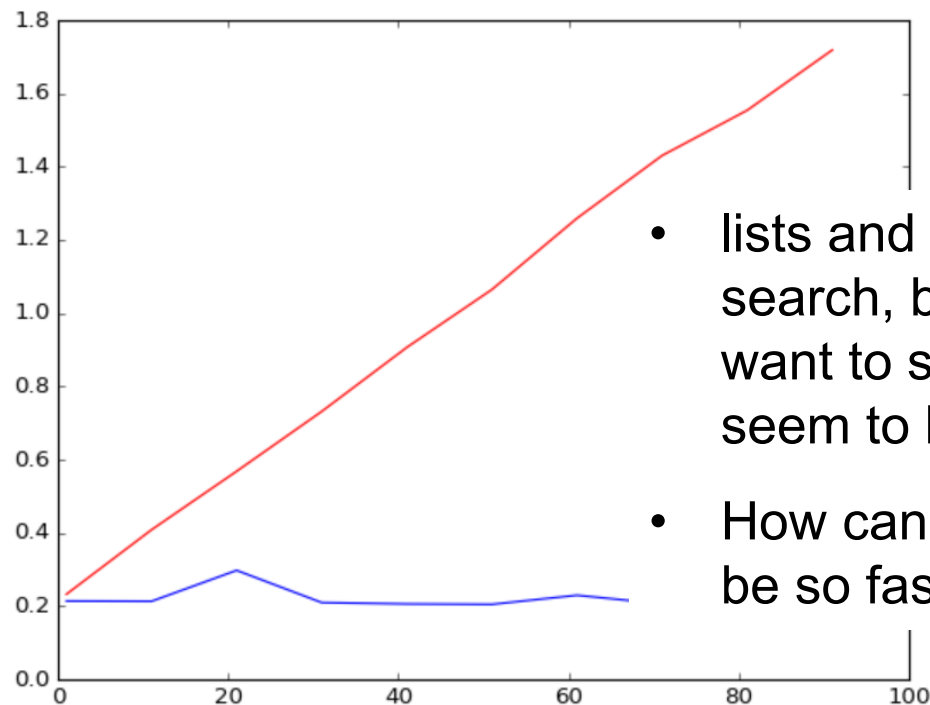
From Lecture 1: searching an element

```
In [5]: timeit.timeit(stmt='1 in A', setup='A = list(range(2, 300))')
```

```
Out[5]: 4.774117301917343
```

```
In [6]: timeit.timeit(stmt='1 in A', setup='A = set(range(2, 300))')
```

```
Out[6]: 0.0499541983165841
```



- lists and sets both allow to search, but if we primarily want to search then sets seem to be the better option
- How can searching on sets be so fast?

Abstract data types and Data Structures

Data Structure

a way to store and organize data to facilitate access and modifications.

Ex. array, hash table, ... later in the course: linked list, heap, ...

Abstract Data Type (ADT)

a set of data values and associated operations that are precisely specified independent of any particular implementation.

Ex. dictionary, ... later in the course: stack, queue, priority queue, ...

- ❑ ADT describe the functionality of data structures
- ❑ Data structures **implement** ADT
 - how is the data stored?
 - which algorithms implement the operations?

Abstract data types and Data Structures

Abstract Data Types

are defined independent of their implementation.

- ❑ We can focus on solving the problem instead of the implementation details
- ❑ Reduce logical errors by preventing direct access to the implementation
- ❑ Implementation can be changed
- ❑ We can have multiple, different implementations for the same data type
- ❑ Easier to manage and divide larger programs into smaller modules

Dictionary

Dictionary

Stores a set S of **elements**, each with an associated **key** (integer value).

Operations

Search(S, k): return a pointer to an element x in S with $\text{key}[x] = k$, or **NIL** if such an element does not exist.

Insert(S, x): inserts element x into S , that is, $S \leftarrow S \cup \{x\}$

Delete(S, x): remove element x from S

S : personal data

- **key**: burger service number
- name, date of birth, address, ... (**satellite data**)

Dictionaries in Python

- ❑ Dictionaries are available in Python as `dict`
- ❑ keys don't need to be integers but for instance can also be strings
- ❑ `dict` is implemented using hash tables, which we will look at in detail today
- ❑ a `set` in Python is like a dictionary but the elements consist of the keys only
- ❑ keys do not need to be integers
- ❑ `set` is also implemented as hash table

Implementing a dictionary

	Search	Insert	Delete
array*	$\Theta(n)$	$\Theta(1)$	$\Theta(1)$
sorted array	$\Theta(\log n)$	$\Theta(n)$	$\Theta(n)$

Today

hash tables

* $\Theta(1)$ Insert and delete for arrays assumes that we have allocated enough (but not more than $O(n)$) memory or dynamically allocate memory

Hash Tables

Hash tables

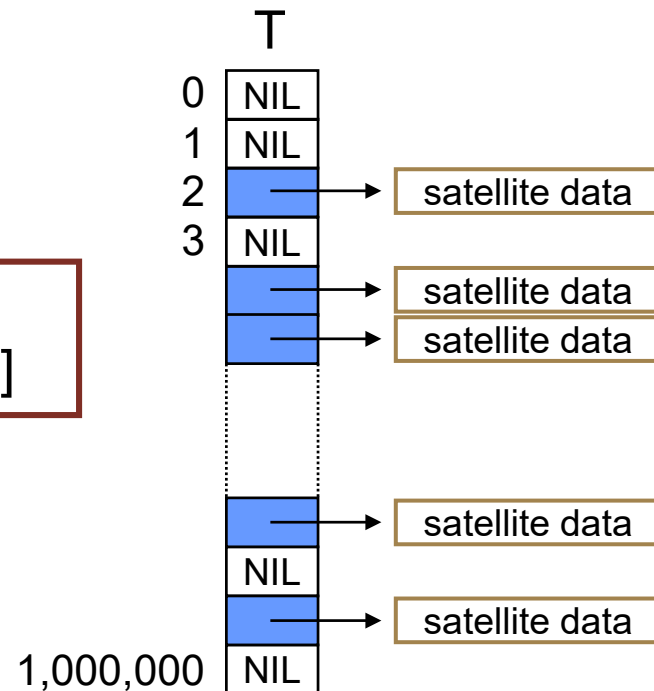
- Hash tables generalize ordinary arrays

Hash tables

- S: personal data in population register
 - key: bsn (burgerservicenummer)
 - name, date of birth, address, ... (satellite data)

Assume: bsn-numbers are integers in the range $[0 \dots 1,000,000]$

Direct addressing
use table $T[0 \dots 1,000,000]$



Direct-address tables


- S : set of elements
 - **key**: unique integer from the universe $U = \{0, \dots, M-1\}$
 - **satellite data**

- use table (array) $T[0..M-1]$

$$T[i] = \begin{cases} \text{NIL} & \text{if there is no element with key } i \text{ in } S \\ \text{pointer to the satellite data} & \text{if there is an element with key } i \text{ in } S \end{cases}$$

Analysis:

- Search, Insert, Delete: $O(1)$
- Space requirements: $O(M)$



Direct-address tables

- S: personal data
 - key: bsn
 - name, date of birth, address, ... (satellite data)
-

Assume: bsn are integers with 9 digits

→ use table $T[0 \dots 999,999,999]$?!?

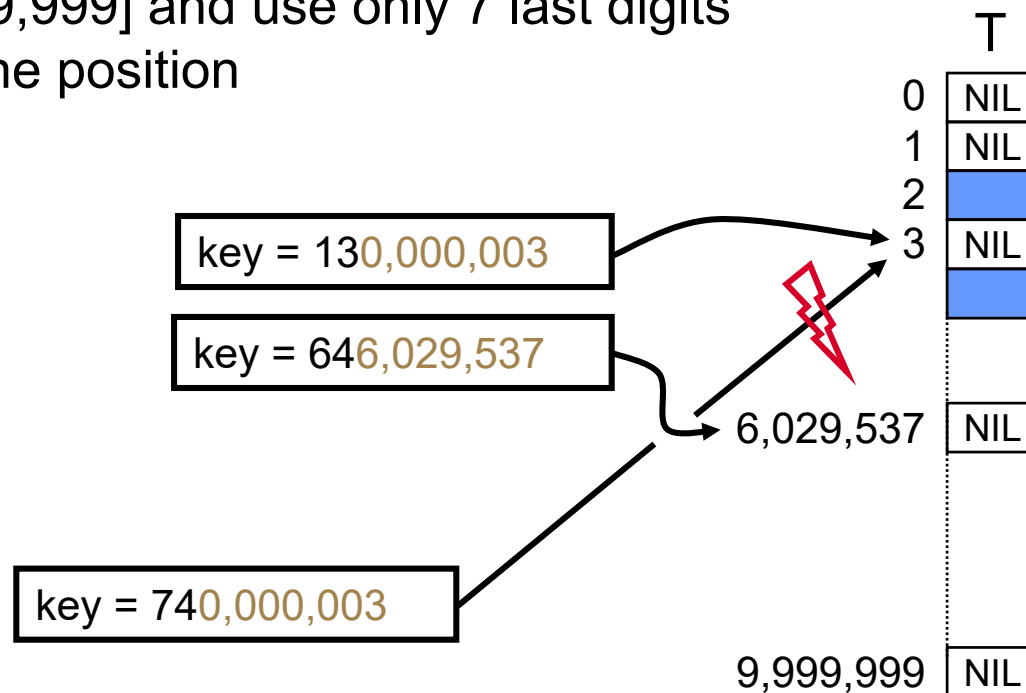
- uses too much memory, most entries will be NIL ...
- if the universe U is large, storing a table of size $|U|$ may be impractical or impossible
- often the set K of keys actually stored is small, compared to U
 - most of the space allocated for T is wasted.

Hash tables

□ S: personal data

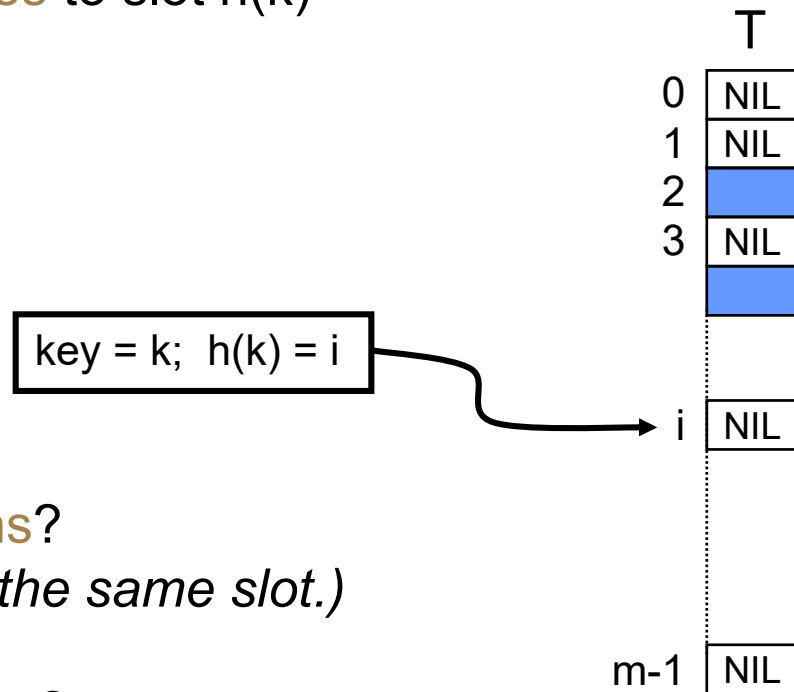
■ key = bsn = integer from $U = \{0 \dots 999,999,999\}$

Idea: use a smaller table, for example,
 $T[0 \dots 9,999,999]$ and use only 7 last digits
to determine position



Hash tables

- S set of keys from the universe $U = \{0 \dots M-1\}$
 - use a hash table $T[0..m-1]$ (with $m \leq M$)
 - use a hash function $h : U \rightarrow \{0 \dots m-1\}$ to determine the position of each key: key k hashes to slot $h(k)$



- How do we resolve collisions?
(Two or more keys hash to the same slot.)
- What is a good hash function?

Resolving collisions: chaining

Chaining: put all elements that hash to the same slot into a linked list

Example ($m=1000$):

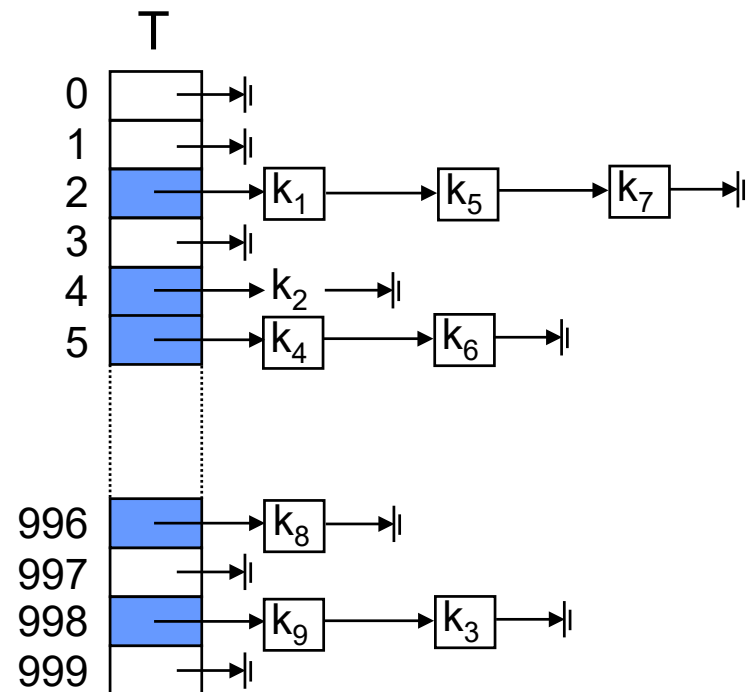
$$h(k_1) = h(k_5) = h(k_7) = 2$$

$$h(k_2) = 4$$

$$h(k_4) = h(k_6) = 5$$

$$h(k_8) = 996$$

$$h(k_9) = h(k_3) = 998$$



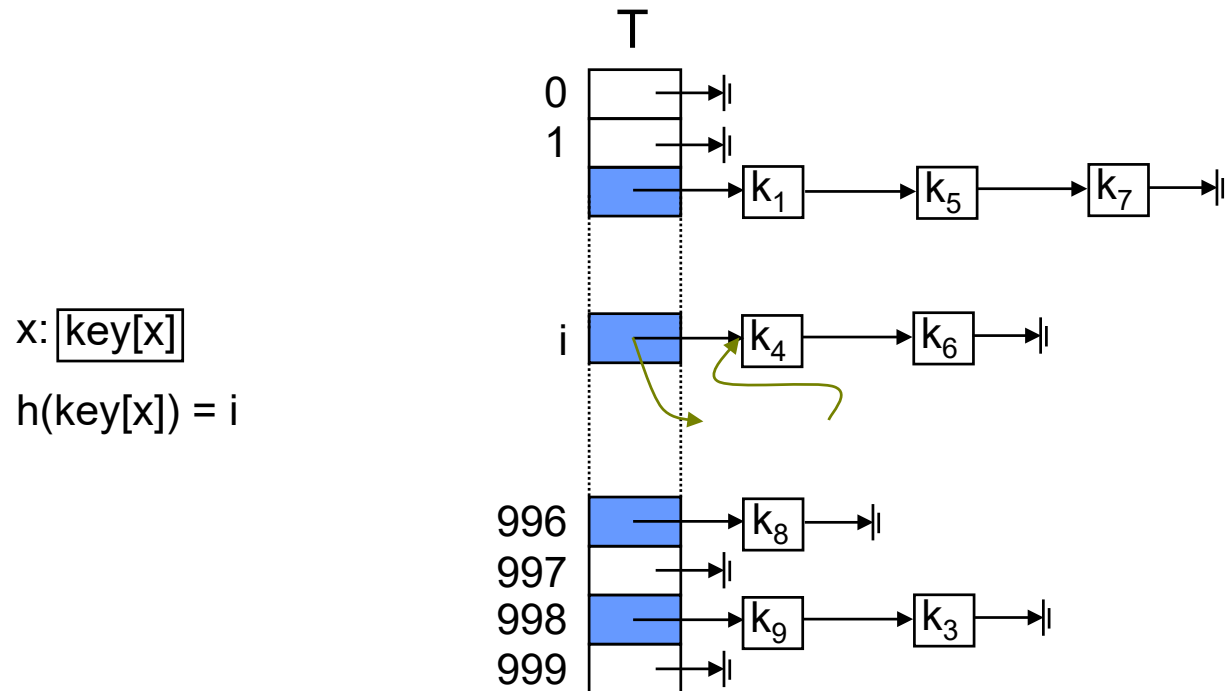
Pointers to the satellite data also need to be included ...

Hashing with chaining: dictionary operations

Chained-Hash-Insert(T, x)

insert x at the head of the list $T[h(\text{key}[x])]$

Time: $O(1)$



Hashing with chaining: dictionary operations

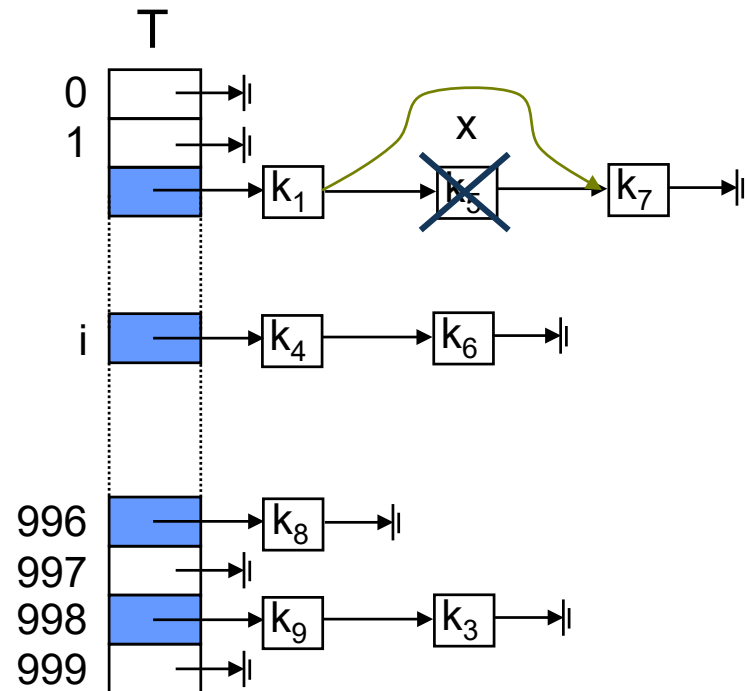
Chained-Hash-Delete(T, x)

delete x from the list $T[h(\text{key}[x])]$

□ x is a pointer to an element

Time: $O(1)$

(store pointers to previous and next element, update these pointers for previous and next element)



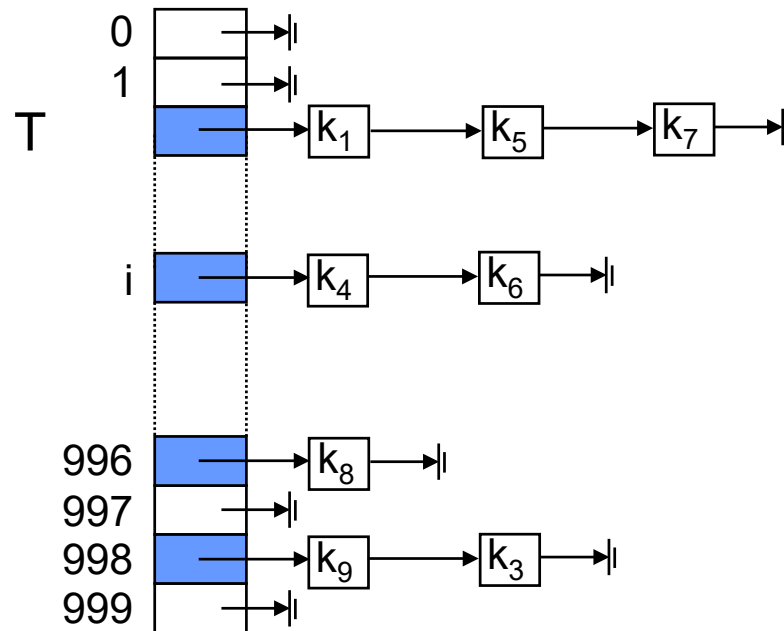
Hashing with chaining: dictionary operations

Chained-Hash-Search(T, k)

search for an element with key k in list $T[h(k)]$

Time:

- unsuccessful: $O(1 + \text{length of } T[h(k)])$
- successful: $O(1 + \# \text{ elements in } T[h(k)] \text{ ahead of } k)$



Hashing with chaining: analysis

Time:

- unsuccessful: $O(1 + \text{length of } T[h(k)])$
- successful: $O(1 + \# \text{ elements in } T[h(k)] \text{ ahead of } k)$

→ worst case $O(n)$

Can we say something about the average case?

Simple uniform hashing

any given element is equally likely to hash into any of the m slots

Hashing with chaining: analysis

Simple uniform hashing

any given element is equally likely to hash into any of the m slots

in other words ...

- the hash function distributes the keys from the universe U uniformly over the m slots
 - the keys in S , and the keys with whom we are searching, behave as if they were randomly chosen from U
- ➔ we can analyze the average time it takes to search as a function of the **load factor** $\alpha = n/m$

(m : size of table, n : total number of elements stored)

Hashing with chaining: analysis

Theorem

In a hash table in which collision are resolved by chaining, an unsuccessful search takes time $\Theta(1+\alpha)$, on the average, under the assumption of simple uniform hashing.

Proof (for an arbitrary key)

- the key we are looking for hashes to each of the m slots with equal probability
 - the average search time corresponds to the average list length
 - average list length = total number of keys / # lists = α
-
- The $\Theta(1+\alpha)$ bound also holds for a successful search (although there is a greater chance that the key is part of a long list).
 - If $m = \Omega(n)$, then a search takes $\Theta(1)$ time on average.

What is a good hash function?

What is a good hash function?

1. as random as possible
get as close as possible to simple uniform hashing ...
 - the hash function distributes the keys from the universe U uniformly over the m slots
 - the hash function has to be as independent as possible from patterns that might occur in the input
2. fast to compute

What is a good hash function?

Example: hashing performed by a compiler for the symbol table

- keys: variable names which consist of (capital and small) letters and numbers: i, i2, i3, Temp1, Temp2, ...

Idea:

- use table of size $(26+26+10)^2$
- hash variable name according to the first two letters:
Temp1 \rightarrow Te




Bad idea: too many “clusters”
(names that start with the same two letters)

What is a good hash function?

Assume: keys are natural numbers

if necessary first map the keys to natural numbers



➔ the hash function is $h: \mathbf{N} \rightarrow \{0, \dots, m-1\}$

▣ the hash function always has to depend on all digits of the input

Common hash functions

Division method: $h(k) = k \bmod m$

Example: $m=1024$, $k = 2058 \rightarrow h(k) = 10$

- don't use a power of 2
 $m = 2^p \rightarrow h(k)$ depends only on the p least significant bits
- use m = prime number, not near any power of two

Multiplication method: $h(k) = \lfloor m (kA \bmod 1) \rfloor$

1. $0 < A < 1$ is a constant
 2. compute kA and extract the fractional part
 3. multiply this value with m and then take the floor of the result
- Advantage: choice of m is not so important, can choose m = power of 2

Resolving collisions

more options ...

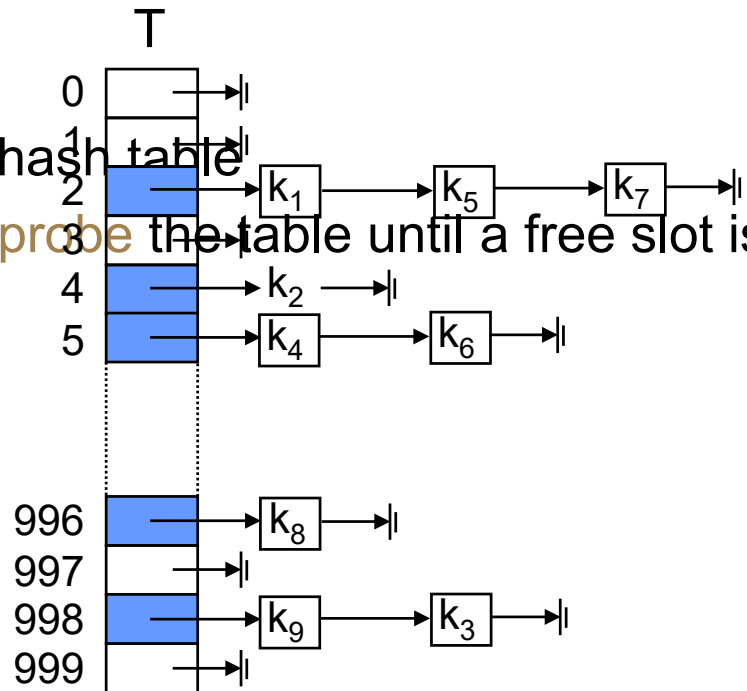
Resolving collisions

Resolving collisions

1. **Chaining**: put all elements that hash to the same slot into a linked list

2. **Open addressing**:

- store all elements in the hash table
- when a collision occurs, **probe** the table until a free slot is found



Hashing with open addressing


Open addressing:

- store all elements in the hash table
- when a collision occurs, **probe** the table until a free slot is found

Example: $T[0..6]$ and $h(k) = k \bmod 7$

1. insert 3
2. insert 18
3. insert 28
4. insert 17

T	
0	28
1	
2	
3	3
4	18
5	17
6	



- no extra storage for pointers necessary
- the hash table can “fill up”
- the load factor α is always ≤ 1

Hashing with open addressing

- there are several variations on open addressing depending on how we search for an open slot
- the hash function has two arguments:
the key and the number of the current probe
→ probe sequence $\langle h(k,0), h(k, 1), \dots h(k, m-1) \rangle$

The probe sequence has to be a permutation of $\langle 0, 1, \dots, m-1 \rangle$ for every key k .

Open addressing: dictionary operations

Hash-Insert(T, k)  we're actually inserting element x with $\text{key}[x] = k$

1. $i = 0$
2. **while** $(i < m)$ **and** $(T[h(k,i)] \neq \text{NIL})$
3. **do** $i = i + 1$
4. **if** $i < m$
5. **then** $T[h(k,i)] = k$
6. **else** “hash table overflow”

Example: Linear Probing

- $T[0..m-1]$
- $h'(k)$ ordinary hash function
- $h(k,i) = (h'(k) + i) \bmod m$
- Hash-Insert($T, 17$)

T		
0	28	
1		
2		
3	3	17
4	18	17
5	17	17
6		

Open addressing: dictionary operations

Hash-Search(T,k)

1. $i = 0$
2. **while** $(i < m)$ **and** $(T[h(k,i)] \neq \text{NIL})$
3. **do if** $T[h(k,i)] = k$
4. **then return** “k is stored in slot $h(k,i)$ ”
5. **else** $i = i + 1$
6. **return** “k is not stored in the table”

Example: Linear Probing

- $h'(k) = k \bmod 7$
 $h(k,i) = (h'(k) + i) \bmod m$
- Hash-Search(T,17)

T	
0	28
1	
2	
3	3
4	18
5	17
6	

17
17
17

Open addressing: dictionary operations

Hash-Search(T,k)

1. $i = 0$
2. **while** $(i < m)$ **and** $(T[h(k,i)] \neq \text{NIL})$
3. **do if** $T[h(k,i)] = k$
4. **then return** “k is stored in slot $h(k,i)$ ”
5. **else** $i = i + 1$
6. **return** “k is not stored in the table”

Example: Linear Probing

- $h'(k) = k \bmod 7$
 $h(k,i) = (h'(k) + i) \bmod m$
- Hash-Search(T,17)
- Hash-Search(T,25)

T		
0	28	
1		
2		
3	3	
4	18	25
5	17	25
6		25

Open addressing: dictionary operations

Hash-Delete(T, k)

1. remove k from its slot
2. mark the slot with the special value DEL

Example: delete 18

T	
0	28
1	
2	
3	3
4	DEL
5	17
6	

- Hash-Search passes over DEL values when searching
- Hash-Insert treats a slot marked DEL as empty
 - ➔ search times no longer depend on load factor
 - ➔ use chaining when keys must be deleted

Open addressing: probe sequences

- $h'(k)$ = ordinary hash function

Linear probing: $h(k,i) = (h'(k) + i) \bmod m$

- $h'(k_1) = h'(k_2) \Rightarrow k_1$ and k_2 have the same probe sequence
 - the initial probe determines the entire sequence
 - \Rightarrow there are only m distinct probe sequences
 - all keys that test the same slot follow the same sequence afterwards
-
- Linear probing suffers from **primary clustering**: long runs of occupied slots build up and tend to get longer
 - \Rightarrow the average search time increases

Open addressing: probe sequences

- $h'(k)$ = ordinary hash function

Quadratic probing: $h(k,i) = (h'(k) + c_1i + c_2i^2) \bmod m$

- $h'(k_1) = h'(k_2) \Rightarrow k_1$ and k_2 have the same probe sequence
 - the initial probe determines the entire sequence
 - \Rightarrow there are only m distinct probe sequences
 - but keys that test the same slot do not necessarily follow the same sequence afterwards
-
- quadratic probing suffers from **secondary clustering**: if two distinct keys have the same h' value, then they have the same probe sequence

Note: c_1 , c_2 , and m have to be chosen carefully, to ensure that the whole table is tested.

Open addressing: probe sequences

□ $h'(k)$ = ordinary hash function

Double hashing: $h(k,i) = (h'(k) + i h''(k)) \bmod m$,
 $h''(k)$ is a second hash function

- keys that test the same slot do not necessarily follow the same sequence afterwards
- h'' must be relatively prime to m to ensure that the whole table is tested.
- $O(m^2)$ different probe sequences

Open addressing: analysis

Uniform hashing

each key is equally likely to have any of the $m!$ permutations of $\langle 0, 1, \dots, m-1 \rangle$ as its probe sequence

Assume: load factor $\alpha = n/m < 1$, no deletions

Theorem

The average number of probes is

- $\Theta(1/(1-\alpha))$ for an unsuccessful search
- $\Theta((1/\alpha) \log (1/(1-\alpha)))$ for a successful search

Open addressing: analysis

Theorem

The average number of probes is

- $\Theta(1/(1-\alpha))$ for an unsuccessful search
- $\Theta((1/\alpha) \log(1/(1-\alpha)))$ for a successful search

Proof: $E[\text{\#probes}] = \sum_{1 \leq i \leq n} i \cdot \Pr[\text{\# probes} = i]$
 $= \sum_{1 \leq i \leq n} \Pr[\text{\# probes} \geq i]$

$$\Pr[\text{\#probes} \geq i] = \frac{n}{m} \cdot \frac{n-1}{m-1} \cdot \frac{n-2}{m-2} \cdot \dots \cdot \frac{n-i+2}{m-i+2}$$
$$\leq \left(\frac{n}{m}\right)^{i-1} = \alpha^{i-1}$$

$$E[\text{\#probes}] \leq \sum_{1 \leq i \leq n} \alpha^{i-1} \leq \sum_{0 \leq i \leq \infty} \alpha^i = \frac{1}{1-\alpha}$$

Check the CLRS book for details!

Hash tables

- Hash tables generalize ordinary arrays
 - map a large universe to a small table
- How do we resolve collisions?
 - Chaining
 - Open addressing: linear and quadratic probing, double hashing
- What is a good hash function?
 - Division method
 - Multiplication method

Implementing a dictionary

	Search	Insert	Delete
array	$\Theta(n)$	$\Theta(1)$	$\Theta(1)$
sorted array	$\Theta(\log n)$	$\Theta(n)$	$\Theta(n)$
hash table	$\Theta(1)$	$\Theta(1)$	$\Theta(1)$

- Running times are average times and assume (simple) uniform hashing and a large enough table (for example, of size $2n$). Also inserting/deleting in the array will require resizing.

Drawbacks of hash tables: operations such as finding the min or the successor of an element are inefficient.